S7 GPT-4o Transfer Runs

Table S7.1 lists every harness we evaluated on GPT-4o, the backbone used for the comparison with prior methods in Sec. 6.

Table S7.1: Cross-transfer to GPT-4o on PrimeVul (860 functions, 430 pairs). Harnesses learned with DeepSeek V4 Flash (Codex CLI) are evaluated on GPT-4o, which took no part in optimization. Pairs gives correctly classified pairs out of 430 (samples without a prediction count as incorrect). (sel.) marks the harness selected on validation. Only the iterations shown were evaluated.

Iter	N	Pairwise	Pairs	Acc	F1	Prec	Rec	CWE acc	Loc F1	Loc hit	Parse
Harness learned with DeepSeek V4 Flash (Codex CLI)											
0 (seed)	860	16.51	71/430	51.7	0.543	0.516	0.574	14.6	15.5	27.9	0
7	860	22.09	95/430	53.5	0.593	0.527	0.677	15.1	11.9	19.6	0
8	859	18.37	79/430	54.0	0.388	0.579	0.291	20.0	16.9	26.4	1
10	860	23.49	101/430	54.2	0.495	0.551	0.449	15.5	10.7	17.1	0
11 (sel.)	860	20.70	89/430	53.0	0.500	0.534	0.470	15.8	15.0	22.3	0
12	860	23.26	100/430	54.8	0.485	0.563	0.426	14.2	18.1	25.7	0
14	860	21.86	94/430	54.4	0.453	0.566	0.377	14.8	13.0	19.8	0