S12 HOSTED INFERENCE PRICES

OpenRouter entries are based on the model card's profile headline rates on `openrouter.ai`. Input/output price always varies by routing, precision, caching, and discounts, but these were the prices when we ran our experimentation. Tinker prices include any listed discount. We link the model names to the sources. We used Opus through the subscription (`claude -p`).

Model	Service	Input	Output
GLM-5.3	OpenRouter	0.5614	1.764
DeepSeek V4 Flash 0731	OpenRouter	0.022	0.32
GLM-4.7 Flash	OpenRouter	0.06	0.40
GPT-OSS-120B	OpenRouter	0.03	0.17
GPT-4o	OpenRouter	2.50	10.00
Inkling-Small (64K)	Tinker	0.58	1.44
Nemotron 3.5 Lightning BF16 (64K)	Tinker	0.195	0.495
Claude Opus 4.6	Anthropic subscription[a]	5.00	25.00
Claude Opus 5	Anthropic subscription[a]	5.00	25.00

[a] Run through a subscription; prices shown are Anthropic's API list prices, for reference.

Table S12.1: Listed inference prices in USD per million tokens.